





# Introduction to Applied Data Science in Python

QBS 101.5

*Simon Stone✱, Jeremy Mikecz ✱, Carly Bobak✧*

*✱Research Data Services, ✧Research Computing*

*Dartmouth College*

# Applied Data Science in Python

## Instructors

### Carly Bobak, PhD

- Biomedical Informatics Specialist
- PhD in Quantitative Biomedical Sciences from Dartmouth College (2021)
- Teaches QBS 181: Data Wrangling
- Published papers on diverse public health issues, including TB, Covid, HIV, cancer, sepsis, smoking cessation, and obesity.
- Has two dogs who tend to crash QBS events



### Jeremy Mikecz, PhD

- Research Data Science Specialist
- PhD in Latin American History, University of California
- Offers consultations and workshops on computational text analysis, data visualization, GIS and mapping, Python & R programming
- Published various articles and a book (in-press) showcasing how visualization and mapping can aid historical research



### Dr. Simon Stone

- Research Data Science Specialist
- Doctoral degree in Electrical and Computer Engineering from Technische Universität Dresden
- Offers consultations and workshops on Data Science, Machine Learning, and Software Development
- Published papers on speech technology, sensor design, signal processing, and machine learning
- Freelance software developer by night



# Applied Data Science in Python

## Course materials

- Brand new class!
- Materials are created as we go along
- You can help shape them!
- Materials will consist of lecture slides and code-along notebooks



### Class repository:

<https://github.com/Simon-Stone/qbs-applied-data-science>

### Get the materials:

```
git clone https://github.com/Simon-Stone/qbs-applied-data-science
```

### Update the materials:

```
git pull
```



# Introduction **Data Science**

## **Challenge:**

## Define *Data Science*!

Data science is the study of data to extract meaningful insights for business.

<https://aws.amazon.com/what-is/data-science/>



Data science [...] uncover[s] actionable insights hidden in an organization's data.

<https://www.ibm.com/topics/data-science>



[Data science] models and analyzes key data to continually improve how businesses utilize data.

<https://cloud.google.com/learn/training/machinelearning-ai>





# Introduction Data Science

Data Science is **OSEMN**!\*



Obtain



Scrub



Explore



Model



iNterpret

Data  
Engineer

- Collect
- Clean

Data  
Analyst

- Clean
- Exploratory Data Analysis
- Build and assess model

Machine  
Learning  
Engineer

- Model implementation
- Deployment

Data Scientist

\*pronounced “awesome” - /'ɔ.səm/  
<https://www.datascience-pm.com/osemn>



# Outline

 Introduction

 Schedule

 Learning Objectives


 What makes a Data Science project special?

 What makes a Python project special?

 Tools of the Trade

 Summary

 Assignment





# Learning Objectives

1. **Complete a data science project** through the process of finding, cleaning, compiling, exploring, analyzing, visualizing, and modeling data.
2. **Become proficient at writing and applying Python** code to complete these data science steps.
3. **Practice effectively communicating** (and providing critical feedback on) data science research.
4. Apply the principles of **reproducible research** throughout the data science lifecycle.





# Format

## Class times and location:

- Monday, 4 pm – 5:30 pm
  - Wednesday, 4 pm – 5:30 pm
  - Remsen 312
- 
- Classes will consist of lectures, code-alongs, hands-on project work, or a mix of all the above
  - There will be assignments to continue the engagement with each week's content
  - At the end of each Wednesday class, there is time for project check-ins, show & tell, or troubleshooting
  - Each student is expected to hand in a final project



# Format

## Final Project

### Project requirements:

- Devise a business or research question that can be answered through data
- Identify a suitable dataset
- Implement all major steps of a data science project to find an answer to the chosen question

### Deliverables:

- Project pitch (details to come)
- Fully reproducible project repository
- Final presentation



# Schedule

## Week 1 (June 26 & June 28):

### Getting Started with Applied Data Science & Python

- Introduction to Applied Data Science with Python
- Data Wrangling

## Week 2 (July 3 & July 5):

### Texts, Maps, and Graphs

- Data Visualization
- NLP I

## Week 3 (July 10 & July 12):

### Texts as Data

- NLP II
- NLP III

## Week 4 (July 17 & July 19):

### Machine Learning

- Scikit-learn
- PyTorch

## Week 5 (July 24 & July 26):

### Project Work

- Project proposals due (Monday July 24)
- *No class meetings*

## Week 6 (July 31 & August 2):

- *No class meetings*

## Week 7 (August 7 & August 9):

### Collaboration and Documentation

- Documenting, Sharing and Evaluating Code
- Peer Review

## Week 8 (August 14 & August 16):

### Code Revision Work

- *No Class Meetings*

## Week 9 (August 21 & August 23):

- Final Presentations and revised code due (Wed. Aug 23)



# Policy on Generative AI

**Use of generative AI is generally accepted in this course if it meets the following requirements:**

- You are expected to take full ownership of your work, including the use of generative AI
- You need to be able to understand and reflect any output of generative AI you want to include in your work
  - “It’s a tool, you’re its master!”
- Just like with code taken from other sources (e.g., Stack Overflow), any code other than your own must be cited (see Week 7: Documentation)
- You are expected to be able to explain and discuss every line of code in your final project

# Applied Data Science with Python

## What makes a Data Science project special

Typical project has **two phases**:

### 1. Exploration phase

 Explore data, imputations, feature engineering, model selection

### 2. Reporting and “deployment” phase

 Generate result tables and graphs, refactor code, export trained models

Subject matter domains differ, tasks are similar

 Multiple stakeholders to consider, “interrogate”, and convince

# Applied Data Science with Python

## **What makes a Python project special**

- 💪 Powerful language with incredible community and industry support
  - Often large number of (inter-)dependencies
- 🚀 Rapid prototyping
- 💻 (Mostly) cross-platform compatible on Windows, Mac OS, Linux
- 🌀 Sprawling, sometimes confusing ecosystem
- 😓 Deployment / reproducibility can be challenging



# Applied Data Science with Python

## **Tools of the Trade**

- Integrated Development Environment (IDE) for writing code
- Command line scripts for project setup and configuration
- Jupyter notebooks for prototyping and reporting
- Version control for code (and sometimes data)

Not part of this course:

- Cloud infrastructure (Amazon Web Services, Google Cloud, Microsoft Azure)



# Tools of the Trade

## Anti-patterns

“A commonly-used process, structure or pattern of action that, despite initially appearing to be an appropriate and effective response to a problem, has more bad consequences than good ones.”

*Gang of Four\**

- Not using an IDE
- Not organizing your project in one place
- Using a too simple or too elaborate project setup
- Overusing Jupyter notebooks (sometimes scripts are better)
- Not managing your Python environment

\* Gamma, E., Helm, R., Johnson, R., Vlissides, J. (1995). *Design patterns: elements of reusable object-oriented software*. Boston, MA, USA: Addison-Wesley Longman Publishing Co., Inc.. ISBN: 0-201-63361-2





## Tools of the Trade

### Recommended Setup for this class

#### Recommended **Python** distribution:

 Vanilla Python 3.8 or newer: <https://www.python.org/>

#### Recommended **IDE**:

 Visual Studio Code: <https://code.visualstudio.com/>

#### Recommended **version control** hosting platform

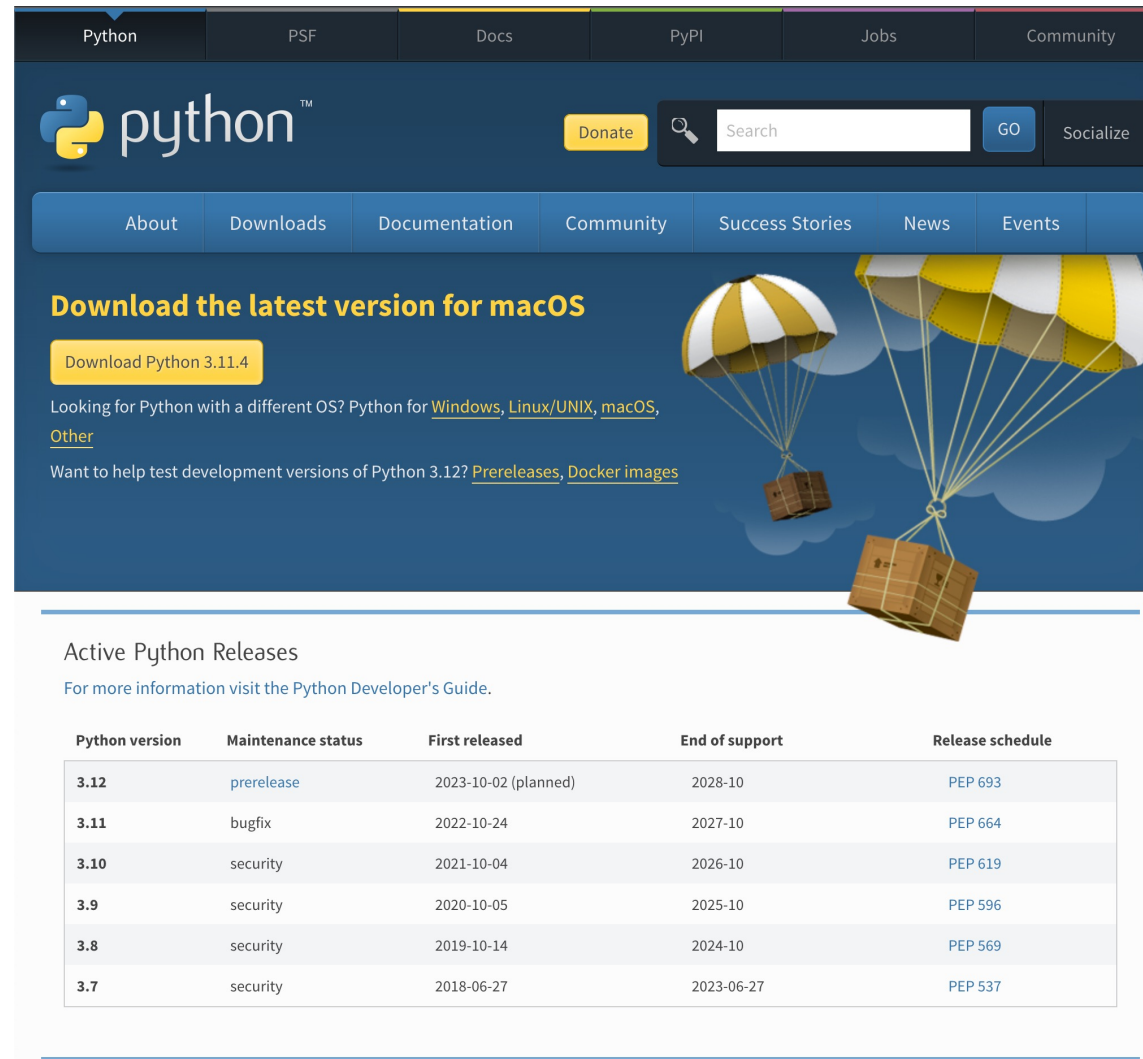
 GitHub: <https://www.github.com>

# Tools of the Trade

## Python

- Recommended minimum **version 3.8+**
- Installer from python.org recommended
- Recommended **environment management**:
  - Using Python's built-in virtualenv:  
<https://docs.python.org/3/library/venv.html>
  - Install packages into virtual environment
  - Document project requirements using:  

```
pip freeze > requirements.txt
```



Download the latest version for macOS

[Download Python 3.11.4](#)

Looking for Python with a different OS? Python for [Windows](#), [Linux/UNIX](#), [macOS](#), [Other](#)

Want to help test development versions of Python 3.12? [Pre-releases](#), [Docker images](#)

### Active Python Releases

For more information visit the [Python Developer's Guide](#).

Python version	Maintenance status	First released	End of support	Release schedule
3.12	prerelease	2023-10-02 (planned)	2028-10	<a href="#">PEP 693</a>
3.11	bugfix	2022-10-24	2027-10	<a href="#">PEP 664</a>
3.10	security	2021-10-04	2026-10	<a href="#">PEP 619</a>
3.9	security	2020-10-05	2025-10	<a href="#">PEP 596</a>
3.8	security	2019-10-14	2024-10	<a href="#">PEP 569</a>
3.7	security	2018-06-27	2023-06-27	<a href="#">PEP 537</a>

<https://www.python.org/downloads/>

# Tools of the Trade

## Python

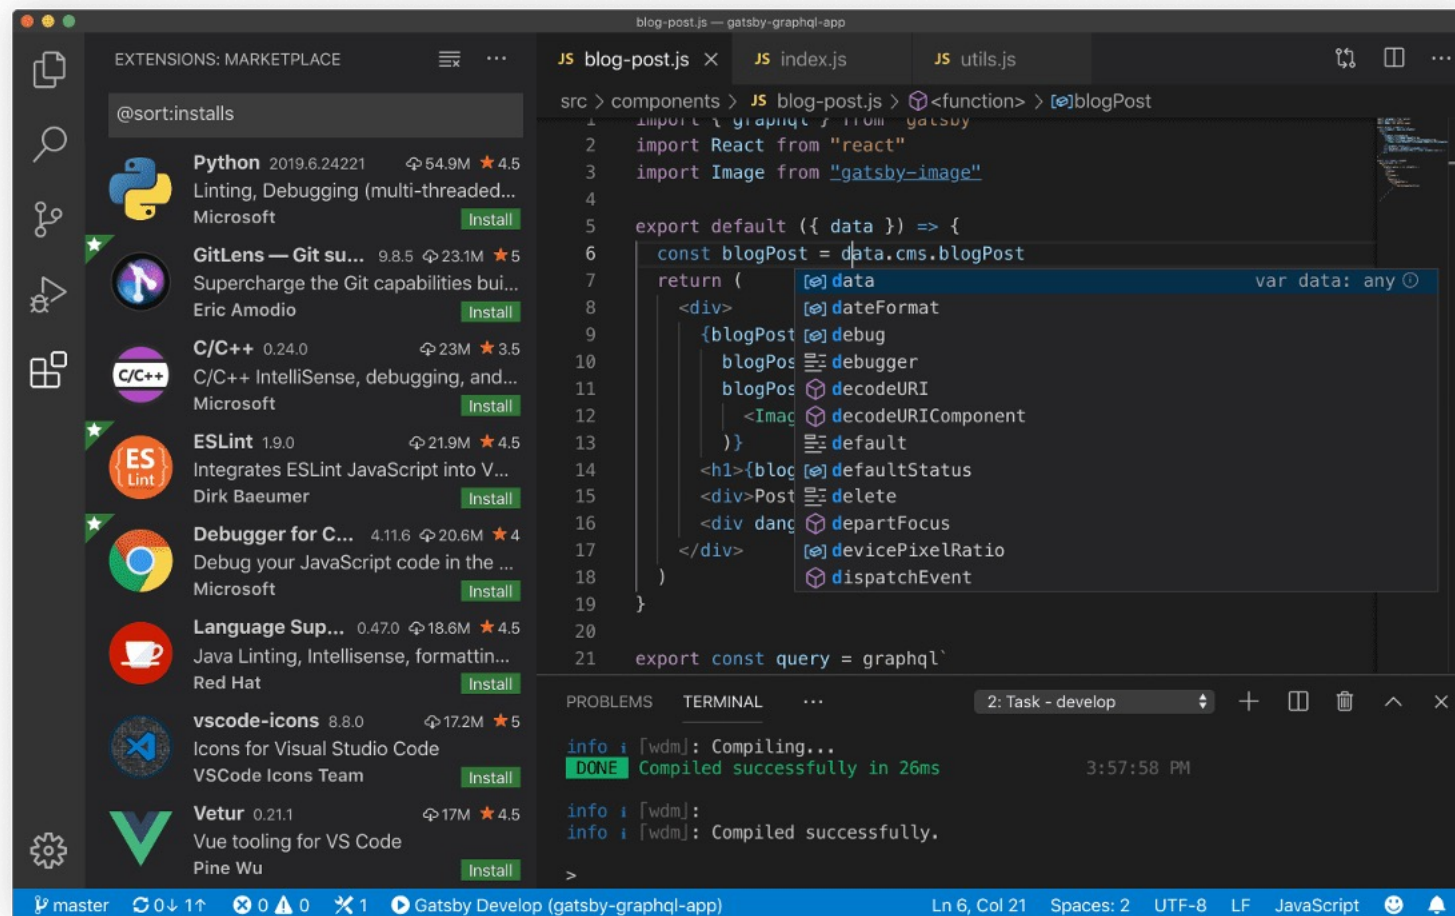
- Python is supported by Jupyter notebooks
- Notebooks are a convenient way to combine code, text, and images in a single file
- Notebooks are great for developing a narrative in your analysis
- Notebooks are not a good choice for project setup, config, or model deployment
- Use regular Python scripts (\*.py) for that



# Tools of the Trade

## Visual Studio Code

- Cross-platform (Windows, Mac, Linux)
- Open source
- Lightweight (no project files required)
- Support for all sorts of languages
- Syntax highlighting and autocomplete
- Extensible in millions of ways
- Natively supports Jupyter notebooks
- Git built-in
- Easy to get started
- Many powerful advanced features



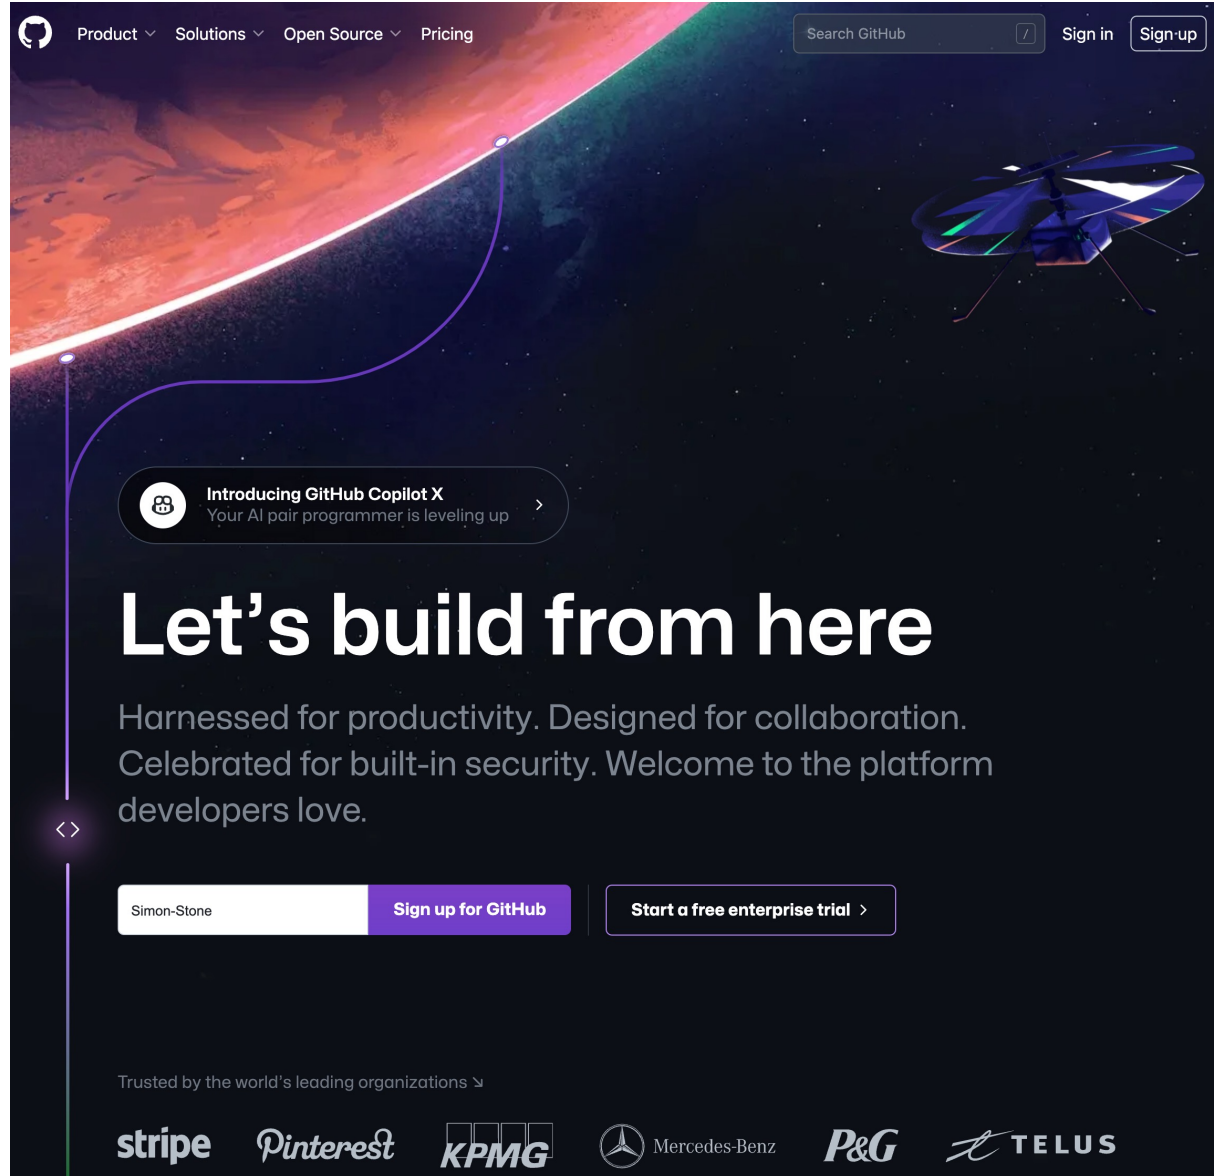
<https://code.visualstudio.com/>



# Tools of the Trade

## GitHub

- Web-based platform for individuals and teams to collaboratively develop software projects
- Uses the Git version control system
- Also offers tools and services for
  - project management,
  - continuous integration and deployment,
  - documentation
- Building up portfolio of projects on GitHub can benefit your job applications



<https://www.github.com/>



# Tools of the Trade

## GitHub

- Repository structure:
  - Many templates available on the internet:
    - <http://drivendata.github.io/cookiecutter-data-science/>
    - <https://gist.github.com/ericmjl/27e50331f24db3e8f957d1fe7bbbe510>
- Try them out, learn what works, adjust and adapt to fit your flow
- Consider creating a GitHub template repo



## Tools of the Trade

### GitHub

- An example structure you can use in this class:
- Tour of the example repo
- Not all of this will be helpful in every project

```
├── data/
│   ├── raw/
│   ├── processed/
│   └── cleaned/
├── models/
├── notebooks/
│   ├── 01-first-logical-notebook.ipynb
│   ├── 02-second-logical-notebook.ipynb
│   ├── prototype-notebook.ipynb
│   └── archive/
│       └── no-longer-useful.ipynb
├── projectname/
│   ├── projectname/
│   │   ├── __init__.py
│   │   ├── config.py
│   │   ├── data.py
│   │   └── utils.py
│   └── setup.py
├── results/
├── scripts/
│   ├── script1.py
│   ├── script2.py
│   └── archive/
│       └── no-longer-useful.py
├── .gitignore
├── README.md
└── requirements.txt
```



# Tools of the Trade

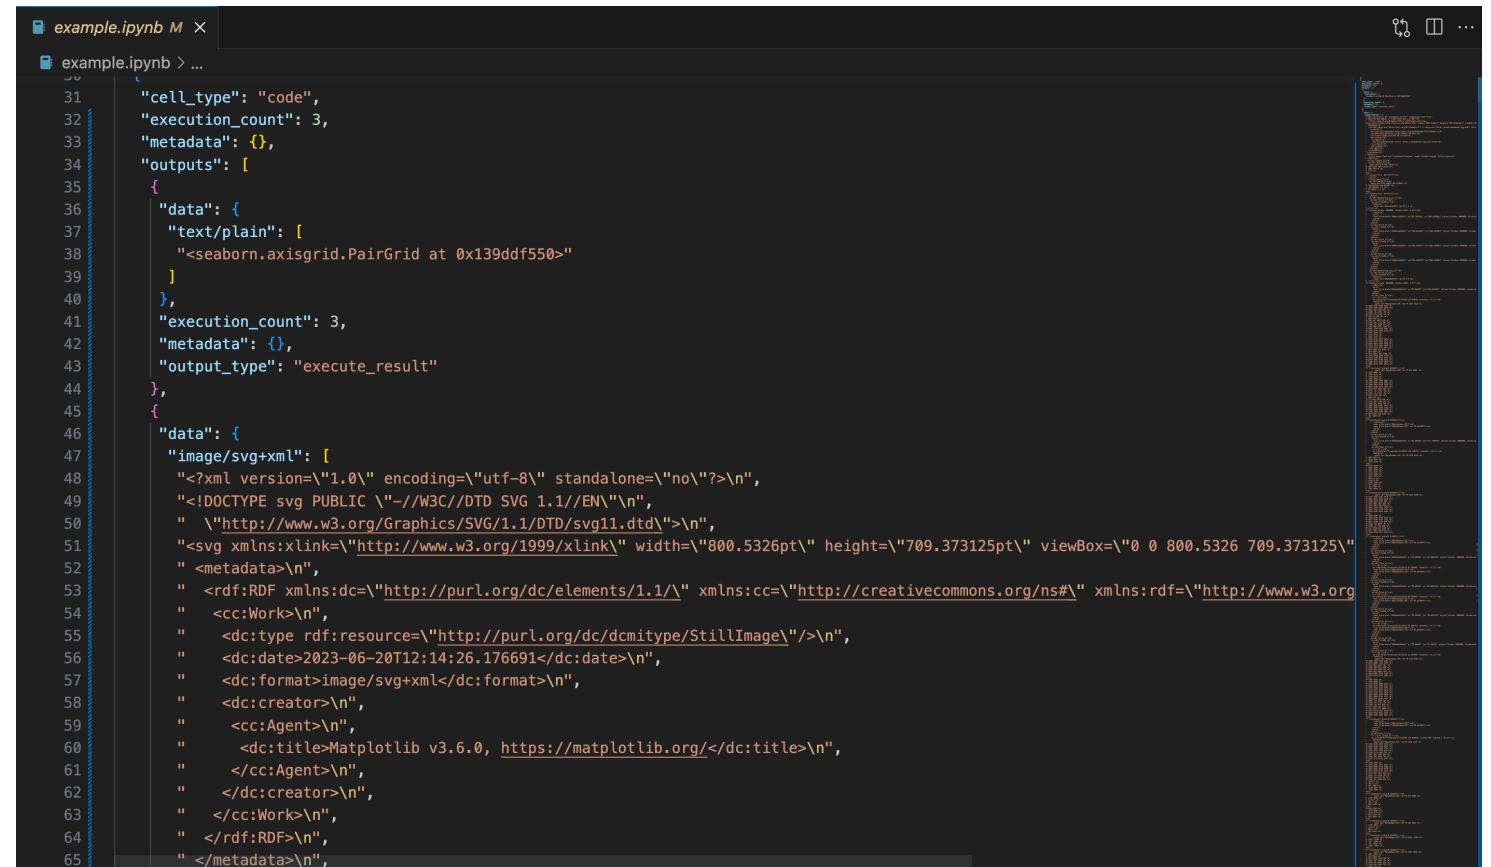
## Bonus tip: Jupyter notebooks in Version Control

### The problem:

- Jupyter notebooks contain metadata, binary blobs, or other artifacts from execution
- Git tracks changed lines in text files
- Diffs become almost meaningless

### Solution:

- Always clear the output before committing
- Use pre-commit hooks
- Use nbstripout  
(<https://github.com/kynan/nbstripout>)



```

example.ipynb M x
example.ipynb > ...
31  "cell_type": "code",
32  "execution_count": 3,
33  "metadata": {},
34  "outputs": [
35    {
36      "data": {
37        "text/plain": [
38          "<seaborn.axisgrid.PairGrid at 0x139ddf550>"
39        ]
40      },
41      "execution_count": 3,
42      "metadata": {},
43      "output_type": "execute_result"
44    },
45    {
46      "data": {
47        "image/svg+xml": [
48          "<?xml version='1.0' encoding='utf-8' standalone='no'?'>\n",
49          "<!DOCTYPE svg PUBLIC \"-//W3C//DTD SVG 1.1//EN\" \"",
50          "  \"http://www.w3.org/Graphics/SVG/1.1/DTD/svg11.dtd\">\n",
51          "<svg xmlns:xlink='\"http://www.w3.org/1999/xlink\"' width='\"800.5326pt\" height='\"709.373125pt\" viewBox='\"0 0 800.5326 709.373125\"'",
52          "  <metadata>\n",
53          "    <rdf:RDF xmlns:dc='\"http://purl.org/dc/elements/1.1/\"' xmlns:cc='\"http://creativecommons.org/ns#\"' xmlns:rdf='\"http://www.w3.org'",
54          "    <cc:Work>\n",
55          "      <dc:type rdf:resource='\"http://purl.org/dc/dcmitype/StillImage\"'/>\n",
56          "      <dc:date>2023-06-20T12:14:26.176691</dc:date>\n",
57          "      <dc:format>image/svg+xml</dc:format>\n",
58          "      <dc:creator>\n",
59          "        <cc:Agent>\n",
60          "          <dc:title>Matplotlib v3.6.0, https://matplotlib.org/</dc:title>\n",
61          "        </cc:Agent>\n",
62          "      </dc:creator>\n",
63          "    </cc:Work>\n",
64          "    </rdf:RDF>\n",
65          "  </metadata>\n",

```





# Assignment

- Set up a skeleton project repository
- Start thinking about a "business/research question" you want to explore in your project
- Start thinking about the kind of data you want to work with
- Start looking for corresponding data
  - Is it already available as a dataset? API? Possible and allowed to scrape?



# Thank you